

Backend Interview Questions

Problem Statement

Grey needs a single backend service that simultaneously:

1. Records every financial event under strict double-entry accounting.

- The system must ensure that for every debit, there is a corresponding credit, and vice versa, following the rules of double-entry accounting.

- **Background on Double-Entry Accounting (DEA):**

- The core equation is:

$$\text{Assets} = \text{Liabilities} + \text{Owner's Equity}$$

- **Assets:** What the bank owns (e.g., cash, buildings, equipment).
 - **Liabilities:** What the bank owes (e.g., users' account balances, salaries).
 - **Owner's Equity:** Contributions the owners made to start or keep the company running (e.g., capital, stock).
- You are free to define services, methods, data flow, and architecture as you see fit. Consider scalability, reliability, and handling of edge cases and failures.

2. Lets users attach any type of file—images, videos, documents, audio—to those events and retrieve them within seconds, anywhere in the world.

- Supported file types:
 - **Images** (e.g., `.jpg`, `.png`, `.heic`)
 - **Videos** (e.g., `.mp4`, `.mov`)
 - **Documents** (e.g., `.pdf`, `.docx`, `.xlsx`)
 - **Audio files** (e.g., `.mp3`, `.wav`)

- Uploaded files should:
 - Be available for **viewing or downloading within seconds** of upload.
 - Be **secure**, with **access control and file validation**.
 - Be **scalable**, supporting **millions of uploads** and **global access**.
 - Optionally support **preview generation** (e.g., image thumbnails, PDF previews).

Scale Assumptions

- The platform must support ~10 million financial transactions and ~50 million file uploads per month.
- Integrity, security, and low latency are non-negotiable.

Your Task

Describe how you would build, deploy, and operate this system end-to-end. Cover:

- Overall architecture and component services
- Data models and storage choices (ledger tables, metadata for files)
- API contracts or interface definitions for posting transactions and uploading/retrieving files
- How you ensure strict debit-equals-credit enforcement, atomicity, and consistency in the ledger
- File ingestion flow (upload URLs or endpoints), validation, storage (hot vs. cold tiers), and global delivery (e.g., CDN design)
- Security controls (authentication, authorization, encryption at rest/in transit, audit logs)
- Scalability strategy (sharding or partitioning, load-balancing, queueing for heavy workloads)
- Reliability and failure handling (idempotency, retries, dead-letter queues, backup/restore, disaster recovery)
- Observability and monitoring (key metrics, alerting strategy)

- Handling of edge cases (e.g., partial failures during a transaction, invalid file formats, corrupted uploads)

Feel free to illustrate your approach with diagrams, pseudo-code, or sequence flows where appropriate. Be explicit about trade-offs and alternative approaches you considered.